



UPPSALA
UNIVERSITET

High Performance Programming **Assignment 4**

Authors:

Amanda Seger

Lovisa Rygaard

Matilda Tiberg

Moa Li Lekander

Uppsala

February 28, 2021

Contents

1	The Problem	1
1.1	The N-body Problem	1
1.2	Time Stepping	1
1.3	The Barnes-Hut Algorithm	1
2	The Solution	3
2.1	Structs	3
2.2	getInput	3
2.3	insert	3
2.4	insertQuadNode	4
2.5	deleteQuadtree	4
2.6	setOutput	4
2.7	calculateNewPositions	4
2.8	calculateForce	5
3	Performance and Discussion	6
3.1	Code Optimizations	6
3.2	Time Complexity	7
3.2.1	Barnes-Hut Efficiency	8
4	Parallelization Using Pthreads	9
Appendix A		11
	Optimized galsim.c from Assignment 4	11
Appendix B		18
	Optimized galsim.c from Assignment 3	18
Appendix C		21
	Parallelized galsim.c from Assignment 3	21
	Parallelized galsim.c from Assignment 4	24

1 The Problem

The goal of the assignment is to implement the Barnes-Hut algorithm in C to reduce the calculations required for simulating the N-body problem. The exact solution to the problem has an $\mathcal{O}(N^2)$ complexity for N amount of particles, ergo, the code quickly becomes slow for an increasing size of the problem. In this assignment, the performance and the complexity of the Barnes-Hut algorithm will be examined and compared to the exact way of solving the N-body problem. Additionally, the two algorithms were parallelized using Pthreads.

1.1 The N-body Problem

The two-body problem consists of determining how two particles, p_i and p_j , interacts with each other depending on their relative position to each other and their masses. According to Newton's law of gravitation, they will exert a force on each other, thus causing the two particles to move. Depending on their relative distance to each other, the forces will differ in both magnitude and direction. Hence, it is a time-dependent system since the forces will cause the particles to move, thus causing the forces to change etc. This system can then easily be expanded to handle N amount of particles, where the total force exerted on particle i is given by equation (1). Here, G is the gravitational constant defined as $100/N$, m_i is the mass of particle i , m_j is the mass of particle j , N is the amount of particles in the system, r_{ij} is the magnitude of the distance between the particles, and \mathbf{r}_{ij} is the distance vector. The constant factor ε_0 is defined as 10^{-5} and exists for the purpose of improving the stability of the system by capping the maximum force between two particles.

$$\mathbf{F}_i = -Gm_i \sum_{j=0, j \neq i}^{N-1} \frac{m_j}{(r_{ij} + \varepsilon_0)^3} \mathbf{r}_{ij} \quad (1)$$

1.2 Time Stepping

The update of positions are done by using the *symplectic Euler* time-integration method, and the governing equations for updating the properties of particle p_i are presented in equation (2). \mathbf{a} denotes the acceleration, \mathbf{u} is the velocity and \mathbf{x} is the position. The current time step is denoted by n , and Δt denotes size of the time steps. This algorithm has a time complexity of $\mathcal{O}(N^2)$, thus causing the time to increase quadratically with an increased amount of particles.

$$\begin{aligned} \mathbf{a}_i^n &= \frac{\mathbf{F}_i^n}{m_i} \\ \mathbf{u}_i^{n+1} &= \mathbf{u}_i^n + \Delta t \mathbf{a}_i^n \\ \mathbf{x}_i^{n+1} &= \mathbf{x}_i^n + \Delta t \mathbf{u}_i^{n+1} \end{aligned} \quad (2)$$

1.3 The Barnes-Hut Algorithm

By using equation (1) to directly sum up the forces acting on particle p_i , the cost of the calculation will significantly grow with an increased number of particles in the simulation. To reduce the calculation time, an approximate method can be used instead. The Barnes-Hut method reduces the time complexity by dividing the domain into subgroups, where many particles can be treated as one object (for details, read below). Then, to calculate the net force on p_i from the particles in said group, only one calculation is performed. This means that if the group contains M particles, then $M - 1$ less calculations of the force needs to be done. The calculations follow equation (1) and (2) in the same manner as before, except that m_j in equation (1) now consists of the central mass of the approximated object instead of a separate particle.

The essential part for this algorithm is the use of a quadtree, which is similar to a binary tree, except a parent node can have four children instead of two. Because of the quadtree, the whole simulation

domain can be stored as the root, and the domain can be divided into four quadrants, where each quadrant is stored in one of the children nodes. Then, if one quadrant contains more than one particle, that quadrant needs to be split into four more quadrants and store its particles in the correct sub-quadrant. The leaf nodes will then contain exactly one particle, whereas the nodes further up in the quadtree will store the center of mass and total mass of all the particles below it. This means that the domain will consist of an adaptive grid with many sub-quadrants, each with either exactly zero or one particles. For an illustrative example, see Figure 1, where the particles are marked as white, and the squares is how the grid was divided. As can be seen, the areas where more particles are situated, the squares are smaller, whereas the squares are bigger at the edge where only a few particles are positioned. The approximation of an object is satisfied if the distance between the particles in the quadrant is small enough compared to the distance between the external particle, p_i , and the midpoint of the quadrant. Mathematically, this is expressed as $h/r \ll 1$, where h is the maximum distance between two particles within the quadrant, and r is the distance to the quadrant.




Figure 1: *Visualization of how the domain can be divided into multiple quadrants.*

In order to achieve an effective algorithm, a way of determining how deep into the tree the traversing should be made must be defined. The value of θ , calculated from equation (3), is used to decide if the node is sufficiently far away from p_i , or if the algorithm should keep traversing down the quadtree. The threshold value is defined as θ_{max} , and if the value of θ exceeds θ_{max} , then the quadtree needs to be traversed further. The threshold can attain the values $\theta_{max} \in [0, 1]$, where a value of $\theta_{max} = 0$ implies that no approximations are made, i.e. the exact solution should be found. Vice versa, a higher value of θ_{max} means that more approximations will be made, thus giving a less accurate result. The value of θ_{max} must therefore be chosen to be as high as possible (to speedup the time required for making the computations) while still achieving the desirable level of accuracy of the calculations.

$$\theta = \frac{\text{width of current quadrant}}{\text{distance between } p_i \text{ and middle of current quadrant}} \quad (3)$$

2 The Solution

The calculations were made in the code displayed in Listing 1 in Appendix A, and all created functions are explained in the sections below. In the code there are also three implemented structs to organize the values used in the functions. The main program is made in such a way that it achieves six extra parameters from the user which are stored in various variables. Moreover, an array of the struct `particles_t` and a struct `quadtree_t` are created to be used in functions to come. The functions `getInput`, `calculateNewPositions` and `setOutput` are returning an integer value to handle possible error messages. If the integer is -1 the main program will first make sure to free all allocated memory before terminating the program. The function `get_wall_seconds` is also included to calculate the run time of the program in wall seconds.

2.1 Structs

The program includes three structs for storing necessary values. The purpose of the struct `particle` is simply to hold all values of a particle. The struct `quadtree` is created to be used as a tree structure holding the necessary values for the Barnes-Hut algorithm. Therefore, it contains the values of the center of mass, a `particle` struct, a `quadrant` struct and four children quadtrees. The center of mass information is the x and y positions and the total mass stored in the struct and its children. The third struct `quadrant` stores the size of the current quadrant in the Barnes-Hut algorithm and the position of the quadrants origin in the whole simulated space. The origin was chosen to be placed in the left bottom corner.

2.2 getInput

The function `int getInput(const char* __restrict filename, quadtree_t** __restrict root, particle_t* __restrict particles, const int N)` uses the given input file `filename` and number of particles `N` to store the input values in the array of particles `particles` and places them in the quadtree `root`. The file is first read with `fread` into an array of doubles. The double array is thereafter used to write all values into their suitable variables in the `particle` array. To add each particle a `quadrant` with the size of the whole space is created and together with the `particle` is used in the function `insert` to update the quadtree `root`.

The function also includes to error checks where an error message is printed and a necessary `free` and `fclose` is used before returning to the main program. At the end of the function the functions `free` and `fclose` are also used.

2.3 insert

The purpose of the function `int insert(quadtree_t** node, particle_t* __restrict particle, quadrant_t* __restrict quadrant)` is to insert a `particle` into a quadtree `node` with the size of the `quadrant`.

There are three possibilities of handling the node, the first case is if the node is `NULL`. In that case, a new `quadtree` is allocated and is given the properties of the input particle together with a new allocated `quadrant` with the same size as the input `quadrant`. All of the children to the node is also set to `NULL`.

Another case is if the input quadtree node is an internal node, meaning that the node itself does not hold information of a single particle, but the children further down the tree structure have particles. In this case, the center of mass of the node, meaning the positions `cm_x` and `cm_y` and the total mass `mass_tot`, is updated to include the new particle to be inserted and afterwards moving the particle further down the tree with the function `insertQuadNode`.

The last case is if the input quadtree node is an external node, which means that the node holds

information of one single particle and all the children nodes are NULL. In this case, the node's already existing particle is given to a new pointer particle so that the particle of the node can be set to NULL and both the old and the new particle can move further down the quadtree to update the children of the node with the function `updateQuadNode`. Afterwards, the center of mass of the node is updated with the new added particle. This case also catches the error of having two particles in the exact same position.

2.4 insertQuadNode

The function `void insertQuadNode(quadtree_t** __restrict node, particle_t* __restrict particle)` finds in which quadrant the position of the input particle lies in order to insert the particle in the correct child of the input node. The first step is that the function creates a new quadrant with half the width and height of the quadrant of the input node, which is the correct size of the child nodes. Afterwards, the x and y positions of the particle are compared to the new width and height to find which new quadrant the particle lies in. When the quadrant is found, the origin of the created quadrant is set accordingly before calling the `insert` function again with the suitable child node, the input particle and the new quadrant. Afterwards, the allocated quadrant is freed.

2.5 deleteQuadtree

To delete a `quadtree` with all its children properly the function `void deleteQuadtree(quadtree_t **node)` is used. The function is recursively working its way down the tree by recalling the function with the child nodes as input. When the input node is zero the return call is used to stop the function, and when all functions have been called the `quadtree` and its assigned `quadrant` are freed and the `quadtree` node is set to NULL.

2.6 setOutput

The function `int setOutput(const char* __restrict filename, particle_t* __restrict data, const int N)` is implemented in a similarly way as `getInput`. The main task of this function is to write the result to the output file. However, the used function `fwrite` can use the created array of particles straight away since all values have been implemented in the correct order as the reference files. Thus, a temporary double array is not necessary to implement in the `setOutput` function.

2.7 calculateNewPositions

To calculate the particle positions over time the function `int calculateNewPositions(quadtree_t** root, particle_t* __restrict particles, const int N, const int nsteps, const double dt, const double theta_max)` was implemented. The function uses the given parameters and also the given value of the gravitational constant `G` to update the position and velocity values of parameters `root` and `particles`. The calculation of the new positions is made with the symplectic Euler time integration method, where the positions and velocities are calculated and updated `nsteps` number of times with the time step `dt`, as given in the input arguments.

Instead of using two loops to calculate the forces on one particle created by the other particles in the system, there is only one loop over `int i` where the function `calculateForce` is called to calculate the force exerted on each particle `i`. The force `F` is a double array and is calculated in the x and y directions, where `F[0]` represents the x direction and `F[1]` represents the y direction. To use the function `calculateForce`, a `quadtree` node is used, but because it is a recursive function which updates the node, the parameter `root` can not be used directly. Instead a new `quadtree` `root_new` is allocated and first given the same values as the input `root`, to be used as the input in `calculateForce`. When the forces have been calculated, they are multiplied with the gravitational constant and the mass

of the particle `i`. All calculations have been made according to equation (1). Afterwards, the accelerations are calculated and used to calculate the new velocities and positions, according to the equations in (2).

If the velocity and position values of the particle in `particle[i]` would be updated immediately it would affect the values of the remaining forces yet to be calculated, and give incorrect results. Therefore, another array of particles `particles_new` was allocated where the new position and velocity values are stored while the old values can continue to be used in calculations. When all calculations are done, the positions and the velocities of the parameter `particle_t* particles` are updated with the values from `particles_new`, so that it can be updated in the main function. Similarly, the new particle values have to be updated in the quadtree `root_new`. This is done by first calling the function `deleteQuadtree` to properly delete all nodes in `root_new` and afterwards inserting all new particles in `root_new` with the function `insert`. The `insert` function also needs an `quadrant` as input, which is allocated at the beginning of `calculateNewPositions` and updated to represent to whole simulated space before calling `insert`.

When all calculations are made, the input `root` is updated and all allocated variables `particles_new`, `root_new`, `F` and `quadrant` are freed. The function also includes two error checks, first if the new calculated positions of the particle exceeds the perimeter of the simulated space and secondly if the insert function returns an error.

2.8 calculateForce

The purpose of the function `void calculateForce(quadtree_t** __restrict node, particle_t* __restrict particle, double* __restrict F, const double theta_max)` is to calculate the force `F` exerted on the parameter `particle` from all other particles in the `quadtree_t` parameter `node`. To reduce the calculations, the value of θ (equation (3)) is calculated for each internal node and if the result is smaller or equal to the threshold value `theta_max`, the center of mass values are used to calculate the force instead of the particle values further down the tree. This is possible because the tree structure is following the Barnes-Hut algorithm.

`calculateForce` is a recursive function, where the first case is if the input node is `NULL` and the return function is called. Secondly, if the input node is an external node, the values of the distance vector `r` and the force in the x (`F[0]`) and y (`F[1]`) directions are calculated. The values of the force are added to the input parameter `F`. The calculations are also using the value of `epsilon` as given in the beginning of the function. However, if the x and y positions of the input particle and the particle of the node are the same, the particles are assumed to be the same and the return call is used instead of the calculations.

It is when the input node is an internal node that the value of θ is calculated in attempt to reduce the calculations. As explained above, if the resulting value is within the threshold, the calculation is treated as a single body problem. However, if θ exceeds the threshold `theta_max`, the calculations continue as a many body problem by recursively calling the function again for each of the four child nodes with the same particle, force vector `F` and threshold `theta_max`.

3 Performance and Discussion

In this section, the method of optimizing the serial code is presented as well as discussed. When optimizing the code, a Macbook Pro, Intel[®] Core[™] i5-6267U CPU @ 2.90GHz together with compiler version *Clang-900.0.39.2* was used. The evaluation was done for the file *ellipse_N_03000.gal* with 200 time steps of 0.00001 seconds each and $\theta_{max} = 0.25$. Additionally, the time complexity is evaluated and discussed. Here, the performance is compared between two computers.

When the optimizations were made, the application Zoom (and possibly a few others) was running in the background, which could have had an impact on the time required for the code to run. However, the final result presented in `best_timing.txt` was obtained when no other applications were running in the background on the computer. To clarify, all optimizations and time measurements in this section were made without parallelization.

3.1 Code Optimizations

The time was measured for different compiler flags, as can be seen in Table 1, where the fastest compiler flag was `-O3 -march=native -ffast-math` and was therefore the one chosen. Then, `-funroll-loops`, the flag for loop unrolling, was added but did not result in either an improvement or deterioration, so it was kept.

Table 1: Performance difference by using different compiler flags in the Makefile. Measured in wall seconds. The quickest run time is marked in bold.

Flags	Time [s]
-O0 (default)	30.104
-O1	7.260
-O2	7.438
-O3	7.280
-Ofast	7.448
-Os	7.435
-O3 -ffast-math	7.477
-O3 -march=native	7.387
-O3 -march=native -ffast-math	7.067
-O2 -march=native -ffast-math	7.125
-Ofast -march=native -ffast-math	7.148

The threshold value, θ_{max} was selected to be as large as possible, to speed up the code, but still keeping the maximum difference between the exact result and the computed result below 10^{-3} . This was done by testing different values of θ_{max} between 0.02 and 0.5 and evaluating the obtained results, and from this $\theta_{max} = 0.25$ turned out to be the optimal threshold value.

In order to reduce the amount of function calls, some optimization changes were tested. In the function `calculateForce`, the function `pow` was replaced with multiplication. This led to a small improvement in the code computation and therefore the change was kept. However, this change negatively affects the readability for the user. Also, the function `fabs` was replaced by the function `sqrt` together with multiplication. This change resulted in a slower code, which was expected since the number of function calls was the same as before and multiplication was needed. That change was therefore removed.

In order to save space, removal of unnecessary variables was tested. In the function `calculateForce`, the two variables for acceleration were removed. This change did not make any improvement on the code regarding time. Since there were only two doubles, meaning that the save in space is very little and removing these would affect the readability in an negative way, the change was not kept.

The keyword `__restrict` was added to the functions where more than one pointer was used as input and where it was guaranteed that they did not point at the same memory space.

"Structure packing" was also added by adding the line `__attribute__((__packed__))` in the declaration of the structs `particle_t` and `quadrant`, in order to minimize the size of the structs used in the code by letting the compiler remove the padding for the structs automatically.

3.2 Time Complexity

In Figure 2, the time consumption plotted as a function of N can be seen for the direct-sum algorithm. In Figure 3 the time consumption for the Barnes-Hut algorithm can be seen. Both algorithms are plotted on a linear scale in these figures, and it can be seen that the direct-sum has a more quadratical appearance, while the plotted result for Barnes-Hut appears to have a more linear dependence (however, not perfectly linear). In Figure 4, the measured time results for the Barnes-Hut method are plotted in logarithmic scale and compared to the theoretical slope of $N \log(N)$, $\log(N)$ and N^2 . This to demonstrate the results visually, and from this plot, it can be concluded that the time complexity for the Barnes-Hut algorithm is $\mathcal{O}(N \log N)$.

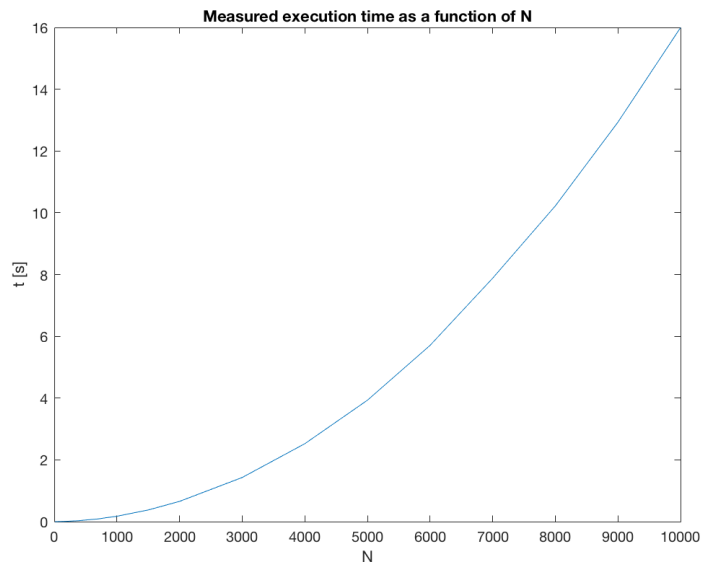


Figure 2: The execution time plotted as a function of N for all available input data, for the $\mathcal{O}(N^2)$ case. $N \in [10, 10000]$.

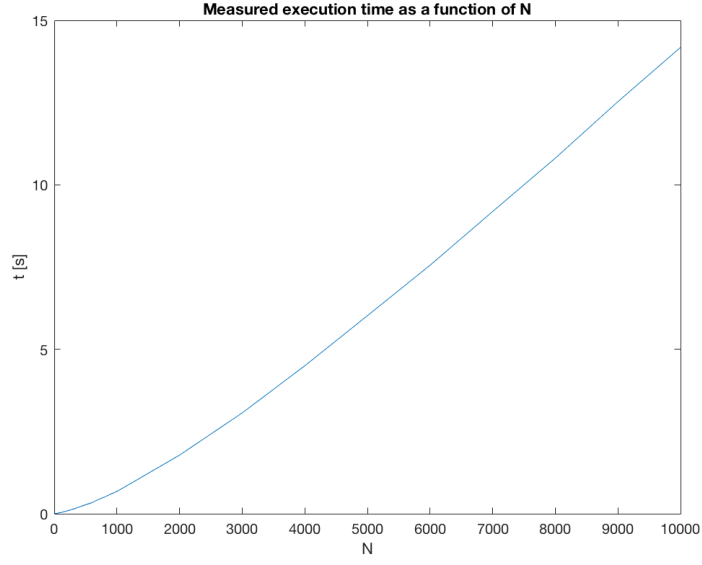


Figure 3: The execution time plotted as a function of N for all available input data, for the Barnes-Hut case. $N \in [10, 10000]$.




Figure 4: Logarithmic plot of the execution time to analyze the time complexity. Here, the slope with logarithmic complexity, $\log(t(N))$ (red), is compared to the slope of the functions: $N\log(N)$ (blue), $\log(N)$ (green) and N^2 (pink).

3.2.1 Barnes-Hut Efficiency

To analyze the efficiency of the implemented Barnes-Hut algorithm compared to the $\mathcal{O}(N^2)$ algorithm, the run time of the programs was measured for different numbers of particles N . Both programs were executed for 200 time steps of 0.00001 seconds. The Barnes-Hut algorithm was made with a threshold of $\theta_{max} = 0.25$.

The measurements were made on two computers with different CPU models and compiler versions. The results are displayed in Figure 2. For the first computer with the CPU model Intel® Core™ i5-3427U and compiler *Apple LLVM version 10.0.0 (clang-1000.10.44.4)* the Barnes-Hut algorithm showed a higher efficiency for 4000 particles and more. However, the computer with the CPU model Intel® Core™ i5-6267U and compiler *Apple LLVM version 9.0.0 (clang-900.0.39.2)* needed 10 000 particles before the Barnes-Hut algorithm was more efficient.

Table 2: Comparison of run times for serial codes using two different algorithms, as well as for two different computers. Compiler versions and input data for the code specified in the paragraph above. The time is measured in wall seconds.

	Intel® Core™ i5-3427U		Intel® Core™ i5-6267U	
N	Direct-sum [s]	Barnes-Hut [s]	Direct-sum [s]	Barnes-Hut [s]
2000	5.467	6.429	1.274	3.434
3000	11.832	11.845	2.961	5.963
4000	20.191	18.548	5.071	8.782
5000	31.559	24.919	7.876	11.746
10000	124.464	66.526	32.132	28.083

4 Parallelization Using Pthreads

In a previous assignment, the direct-sum algorithm for solving the same N-body problem was implemented. That algorithm had efficiency $\mathcal{O}(N^2)$ and the code for that implementation can be found in Appendix B. For the purpose of optimizing the mentioned code as well as the Barnes-Hut algorithm, developed in this assignment, even further, these implementations have been parallelized using Pthreads. The codes for these threaded programs can be found in Appendix C.

In order to parallelize the codes as efficiently as possible, an analysis of what parts of the codes that took the most time to run was made. By using the tool **gprof** it was found that the function `CalculateNewPositions` took up a large majority of the time required for the $\mathcal{O}(N^2)$ algorithm to run, and thus the Pthreads was implemented in that part of the program. To do this, every thread was assigned a certain amount of the particles of which to find the new positions of, and hence a large amount of the computations made in the function `CalculateNewPositions` could be done in parallel.

In the code for the Barnes hut algorithm it was found using **gprof** that most time was spent in the `calculateForce` function. Thus, these computations made by the `calculateForce` was parallelized using Phtreads. Similarly to the parallelization of the $\mathcal{O}(N^2)$ algorithm, each thread was assigned a certain amount of the particles of which to find the new positions of. Thus, much of the computations made by the function `calculateForce` could be done in parallel, which made the code more efficient.

When testing the performance of these threaded programs one of the university's Linux computers was used, running on an Intel® Xeon® E5520 @ 2.27GHz, with a total of 16 CPU:s (4 cores \times 2 sockets, 2 threads per core), and compiler *gcc version 7.5.0 (Ubuntu 7.5.0-3ubuntu1 18.04)*. This was done to get more data points to evaluate the performance's dependence of the available threads. The codes were run with input file *ellipse_N_05000.gal* (which consists of 5000 particles), and 200 time steps of size 0.00001, together with $\theta_{max} = 0.25$.

In Figure 5 the time required for the code to run is shown as a function of the number of threads, for the two algorithms. It can be noted that the time reduces when a larger number of cores is used, for both algorithms. The curves correspond well to the expected result, since the time reduces in a $1/x$ like fashion, where x is the number of threads. When analyzing the time for when 16 threads was used however,

a small increase in the time can be noted when 7 threads were used. This result can be due to the fact that another person seemed to be using the same computer when running the simulations, thus occupying the last available core. It can also be observed that for the relatively large number of particles ($N = 5000$) the performance of the Barnes hut algorithm is significantly better than the $\mathcal{O}(N^2)$ algorithm.

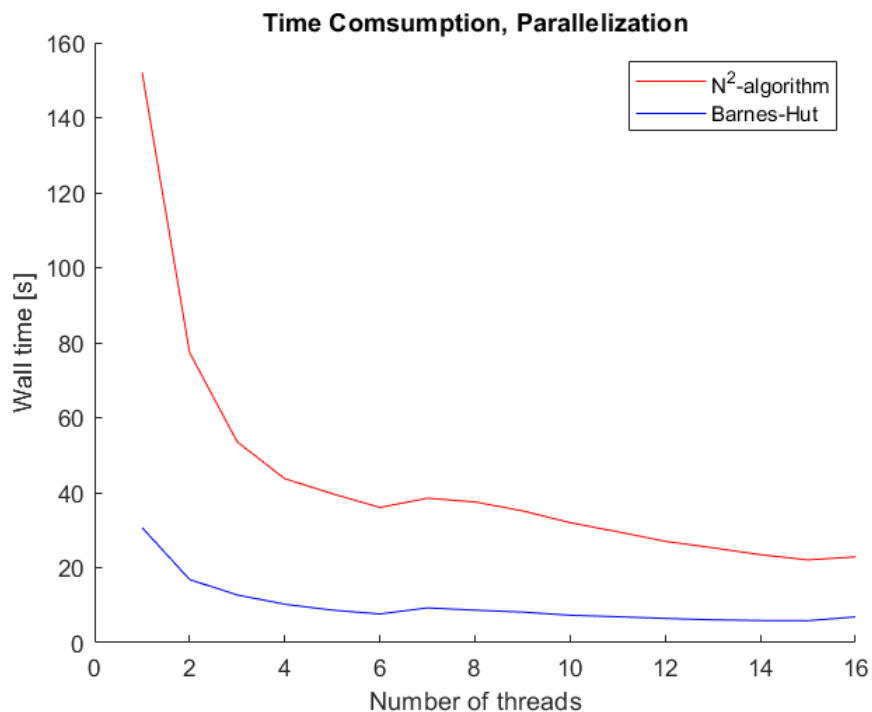


Figure 5: Time required for the two parallelized algorithms to run depending on the number of threads, where the amount of threads range from 1-16. System information can be found in the paragraphs above. The Barnes-Hut algorithm has the complexity $N\log(N)$.

Appendix A

Listing 1: *Optimized galsim.c from Assignment 4*

```
1  #include <stdlib.h>
2  #include <math.h>
3  #include <stdio.h>
4  #include <sys/time.h>
5
6  typedef struct particle
7  {
8      double pos_x;
9      double pos_y;
10     double mass;
11     double vel_x;
12     double vel_y;
13     double brightness;
14 } __attribute__((__packed__)) particle_t;
15
16 typedef struct quadrant
17 {
18     /* Position of specific quadrant in the whole space */
19     double origin_x;
20     double origin_y;
21
22     double width;
23     double height;
24 } __attribute__((__packed__)) quadrant_t;
25
26 typedef struct quadtree
27 {
28     /* center of mass */
29     double mass_tot;
30     double cm_x;
31     double cm_y;
32     particle_t *particle;
33     quadrant_t *quadrant;
34     /* tree nodes */
35     struct quadtree *top_left;
36     struct quadtree *top_right;
37     struct quadtree *bottom_left;
38     struct quadtree *bottom_right;
39 } quadtree_t;
40
41 int getInput(const char* __restrict filename, quadtree_t** __restrict root, particle_t* ...
    __restrict particles, const int N);
42 int insert(quadtree_t** __restrict node, particle_t* __restrict particle, quadrant_t* ...
    __restrict quadrant);
43 void insertQuadNode(quadtree_t** node, particle_t* __restrict particle);
44 void deleteQuadtree(quadtree_t **node);
45 void resetQuadtree(quadtree_t **node);
46 int setOutput(const char* __restrict filename, particle_t* __restrict data, const int N);
47 int calculateNewPositions(quadtree_t** __restrict root, particle_t* __restrict particles, ...
    const int N, const int nsteps, const double dt, const double theta_max);
48 void calculateForce(quadtree_t** __restrict node, particle_t* __restrict particle, double* ...
    __restrict F, const double theta_max);
49
50 /*
51 static double get_wall_seconds() {
52     struct timeval tv;
53     gettimeofday(&tv, NULL);
54     double seconds = tv.tv_sec + (double)tv.tv_usec / 1000000;
55     return seconds;
```

```

56 */
57
58 int main(int argc, char* argv[])
59 {
60     //double time1 = get_wall_seconds();
61     if (argc != 7)
62     {
63         printf("Check your input..\n");
64         return -1;
65     }
66     const int N = atoi(argv[1]);
67     const char* filename = argv[2];
68     const int nsteps = atoi(argv[3]);
69     const double dt = atof(argv[4]);
70     const double theta_max = atof(argv[5]); // theta_max = 0.25
71     //const int graphics = atoi(argv[6]);
72     const char* filename_out = "result.gal";
73     particle_t* particles = (particle_t*)malloc(N*sizeof(particle_t));
74     quadtree_t* root = NULL;
75
76     int err = getInput(filename, &root, particles, N);
77     if(err == -1)
78     {
79         deleteQuadtree(&root);
80         free(particles);
81         return -1;
82     }
83     err = calculateNewPositions(&root, particles, N, nsteps, dt, theta_max);
84     if(err == -1)
85     {
86         deleteQuadtree(&root);
87         free(particles);
88         return -1;
89     }
90
91     err = setOutput(filename_out, particles, N);
92     if(err == -1)
93     {
94         deleteQuadtree(&root);
95         free(particles);
96         return -1;
97     }
98
99     deleteQuadtree(&root);
100    free(particles);
101    //printf("galsim main took %.3f wall seconds.\n", get_wall_seconds()-time1);
102    return 0;
103 }
104
105 int getInput(const char* __restrict filename, quadtree_t* __restrict root, particle_t* ...
106             __restrict particles, const int N)
107 {
108     FILE *stream_in;
109     stream_in = fopen(filename, "rb");
110     double* arr = (double*)malloc(N*6*sizeof(double));
111     quadrant_t *quadrant = (quadrant_t*)malloc(sizeof(particle_t));
112
113     if(stream_in == NULL)
114     {
115         printf("Error: unable to open file: %s\n", filename);
116         fclose(stream_in);
117         free(arr);
118         free(quadrant);
119         return -1;
120     }

```

```

120     size_t input_size = N*6*sizeof(double);
121
122     /* Read input to data array */
123     size_t items_read = fread(arr, sizeof(char), input_size, stream_in);
124     int j = 0;
125     for (int i = 0; i < N*6; i += 6)
126     {
127         particles[j].pos_x = arr[i];
128         particles[j].pos_y = arr[i+1];
129         particles[j].mass = arr[i+2];
130         particles[j].vel_x = arr[i+3];
131         particles[j].vel_y = arr[i+4];
132         particles[j].brightness = arr[i+5];
133
134         quadrant->width = 1;
135         quadrant->height = 1;
136         quadrant->origin_x = 0;
137         quadrant->origin_y = 0;
138
139         insert(root, &particles[j], quadrant);
140         j++;
141     }
142
143     if (items_read != input_size)
144     {
145         printf("Error reading the input file.\n");
146         fclose(stream_in);
147         free(arr);
148         free(quadrant);
149         return -1;
150     }
151     free(quadrant);
152     fclose(stream_in);
153     free(arr);
154     return 1;
155 }
156
157 /* Note: we define origin in bottom-left corner! */
158 int insert(quadtrees_t** node, particle_t* __restrict particle, quadrant_t* __restrict ...
159           quadrant)
160 {
161     if(*node == NULL)
162     {
163         quadtrees_t * new_node = (quadtrees_t*)malloc(sizeof(quadtrees_t));
164         quadrant_t * new_quadrant = (quadrant_t*)malloc(sizeof(quadrant_t));
165         new_node->particle = particle;
166         new_node->mass_tot = particle->mass;
167         new_node->cm_x = particle->pos_x;
168         new_node->cm_y = particle->pos_y;
169
170         new_quadrant->width = quadrant->width;
171         new_quadrant->height = quadrant->height;
172         new_quadrant->origin_x = quadrant->origin_x;
173         new_quadrant->origin_y = quadrant->origin_y;
174         new_node->quadrant = new_quadrant;
175
176         new_node->top_left = NULL;
177         new_node->top_right = NULL;
178         new_node->bottom_left = NULL;
179         new_node->bottom_right = NULL;
180         *node = new_node;
181     }
182     /*
183     * internal node: a node which already contains other nodes

```

```

184     * external node: a node which is initially empty - after its been filled it will only ...
        contain one particle
185     */
186     else if ((*node)->particle == NULL) // internal node
187     {
188         /* updating center of mass */
189         (*node)->cm_x = ((*node)->cm_x*(*node)->mass_tot + ...
            particle->pos_x*particle->mass)/((*node)->mass_tot + particle->mass);
190         (*node)->cm_y = ((*node)->cm_y*(*node)->mass_tot + ...
            particle->pos_y*particle->mass)/((*node)->mass_tot + particle->mass);
191         (*node)->mass_tot += particle->mass;
192         insertQuadNode(&(*node), particle);
193     }
194     else // external node: move the new and old nodes into a new lower quadrant
195     {
196         if((*node)->particle->pos_x == particle->pos_x && (*node)->particle->pos_y == ...
            particle->pos_y){
197             printf("Error: Particles at the same position \n");
198             return -1;
199         }
200         /* Making our node internal by removing the particle */
201         particle_t *particle_old = (*node)->particle;
202         (*node)->particle = NULL;
203         insertQuadNode(&(*node), particle);
204         insertQuadNode(&(*node), particle_old);
205
206         /* updating center of mass */
207         (*node)->cm_x = ((*node)->cm_x*(*node)->mass_tot + ...
            particle->pos_x*particle->mass)/((*node)->mass_tot + particle->mass);
208         (*node)->cm_y = ((*node)->cm_y*(*node)->mass_tot + ...
            particle->pos_y*particle->mass)/((*node)->mass_tot + particle->mass);
209         (*node)->mass_tot += particle->mass;
210     }
211     return 0;
212 }
213
214 void insertQuadNode(quadtrees_t** __restrict node, particle_t* __restrict particle)
215 {
216     quadrant_t *quadrant_new = (quadrant_t*)malloc(sizeof(quadrant_t));
217     quadrant_new->width = (*node)->quadrant->width/2.0;
218     quadrant_new->height = (*node)->quadrant->height/2.0;
219
220     if(particle->pos_x < ((*node)->quadrant->origin_x + quadrant_new->width)) // left quadrant
221     {
222         if (particle->pos_y < ((*node)->quadrant->origin_y + quadrant_new->height)) // ...
            bottom left quadrant
223         {
224             quadrant_new->origin_y = (*node)->quadrant->origin_y;
225             quadrant_new->origin_x = (*node)->quadrant->origin_x;
226             insert(&(*node)->bottom_left, particle, quadrant_new);
227         }
228         else // top left quadrant
229         {
230             quadrant_new->origin_y = (*node)->quadrant->origin_y + quadrant_new->height;
231             quadrant_new->origin_x = (*node)->quadrant->origin_x;
232             insert(&(*node)->top_left, particle, quadrant_new);
233         }
234     }
235     else { // right quadrant
236         if (particle->pos_y < ((*node)->quadrant->origin_y + quadrant_new->height)) // ...
            bottom right quadrant
237         {
238             quadrant_new->origin_x = (*node)->quadrant->origin_x + quadrant_new->width;
239             quadrant_new->origin_y = (*node)->quadrant->origin_y;
240             insert(&(*node)->bottom_right, particle, quadrant_new);

```



```

241     }
242     else // top right quadrant
243     {
244         quadrant_new->origin_x = (*node)->quadrant->origin_x + quadrant_new->width;
245         quadrant_new->origin_y = (*node)->quadrant->origin_y + quadrant_new->height;
246         insert(&(*node)->top_right, particle, quadrant_new);
247     }
248     }
249     free(quadrant_new);
250 }
251
252 void deleteQuadtree(quadtrees_t **node)
253 {
254     if(*node == NULL){ // if empty node
255         return;
256     }
257
258     deleteQuadtree(&(*node)->top_left);
259     deleteQuadtree(&(*node)->top_right);
260     deleteQuadtree(&(*node)->bottom_left);
261     deleteQuadtree(&(*node)->bottom_right);
262
263     free((*node)->quadrant);
264     free(*node);
265     *node = NULL;
266 }
267
268 int setOutput(const char* __restrict filename, particle_t* __restrict data, const int N)
269 {
270     FILE *stream_out;
271     stream_out = fopen(filename, "wb");
272     if(stream_out == NULL)
273     {
274         printf("Error: unable to open file: %s\n", filename);
275         fclose(stream_out);
276         return -1;
277     }
278     fwrite(data, sizeof(particle_t), N, stream_out);
279     fclose(stream_out);
280     return 0;
281 }
282
283 int calculateNewPositions(quadtrees_t** root, particle_t* __restrict particles, const int ...
    N, const int nsteps, const double dt, const double theta_max)
284 {
285     int err = 0;
286     double ax, ay;
287     const double G = 100.0/(double)(N);
288     double* F = (double*)malloc(2*sizeof(double));
289     particle_t *particles_new = (particle_t*)malloc(N*sizeof(particle_t));
290     quadrant_t *quadrant = (quadrant_t*)malloc(sizeof(quadrant_t));
291     quadtrees_t **root_new = (quadtrees_t**)malloc(sizeof(quadtrees_t));
292     *root_new = *root;
293
294     for (int n = 0; n < nsteps; n++)
295     {
296         for(int i = 0; i < N; i++)
297         {
298             F[0] = 0;
299             F[1] = 0;
300
301             calculateForce(root_new, &particles[i], F, theta_max);
302
303             F[0] = -G*particles[i].mass*F[0];
304             F[1] = -G*particles[i].mass*F[1];

```

```

305
306     /* Acceleration */
307     ax = F[0]/(particles[i].mass);
308     ay = F[1]/(particles[i].mass);
309     /* Updating velocities */
310     particles_new[i].vel_x = particles[i].vel_x + dt*ax;
311     particles_new[i].vel_y = particles[i].vel_y + dt*ay;
312     /* Updating positions */
313     particles_new[i].pos_x = particles[i].pos_x + dt* particles_new[i].vel_x;
314     particles_new[i].pos_y = particles[i].pos_y + dt* particles_new[i].vel_y;
315     if(particles_new[i].pos_x >1 || particles_new[i].pos_x <0 || particles_new[i].pos_y ...
        >1 || particles_new[i].pos_y <0){
316         printf("Error: Particle outside interval \n");
317         free(particles_new);
318         free(F);
319         free(root_new);
320         free(quadrant);
321         return -1;
322     }
323 }
324 deleteQuadtree(root_new);
325
326 for (int i = 0; i < N; i++)
327 {
328     particles[i].pos_x = particles_new[i].pos_x;
329     particles[i].pos_y = particles_new[i].pos_y;
330     particles[i].vel_x = particles_new[i].vel_x;
331     particles[i].vel_y = particles_new[i].vel_y;
332
333     quadrant->width = 1;
334     quadrant->height = 1;
335     quadrant->origin_x = 0;
336     quadrant->origin_y = 0;
337     err += insert(root_new, &particles[i], quadrant);
338 }
339 if(err!=0)
340 {
341     free(particles_new);
342     free(F);
343     free(root_new);
344     free(quadrant);
345     return -1;
346 }
347 }
348 *root = *root_new;
349 free(particles_new);
350 free(root_new);
351 free(F);
352 free(quadrant);
353 return 0;
354 }
355
356 void calculateForce(quadtree_t** __restrict node, particle_t* __restrict particle, double* ...
    __restrict F, const double theta_max)
357 {
358     double r, theta;
359     const double epsilon = 0.001;
360     /* F[0] = Fx, F[1] = Fy */
361     if ((*node) == NULL)
362     {
363         return;
364     }
365     /* External node */
366     else if ((*node)->particle != NULL)
367     {

```

```

368  /* Same particle */
369  if ((fabs((*node)->particle->pos_x - particle->pos_x) < 0.0000001) && ...
      (fabs((*node)->particle->pos_y - particle->pos_y) < 0.0000001))
370  {
371      return;
372  }
373  /* x_i: particle, x_j: node->particle */
374  r = sqrt((particle->pos_x - (*node)->particle->pos_x)*(particle->pos_x - ...
      (*node)->particle->pos_x) + (particle->pos_y - ...
      (*node)->particle->pos_y)*(particle->pos_y - (*node)->particle->pos_y));
375  F[0] += (*node)->particle->mass*(particle->pos_x - ...
      (*node)->particle->pos_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
376  F[1] += (*node)->particle->mass*(particle->pos_y - ...
      (*node)->particle->pos_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
377  }
378  /* Internal node */
379  else
380  {
381      /* Calculate ratio */
382      theta = (*node)->quadrant->width/sqrt((particle->pos_x - ((*node)->quadrant->origin_x ...
          + (*node)->quadrant->width/2))*(particle->pos_x - ((*node)->quadrant->origin_x + ...
          (*node)->quadrant->width/2)) + (particle->pos_y - ((*node)->quadrant->origin_y + ...
          (*node)->quadrant->height/2))*(particle->pos_y - ((*node)->quadrant->origin_y + ...
          (*node)->quadrant->height/2)));
383
384      /* Single body problem */
385      if (theta ≤ theta_max)
386      {
387          /* x_i: particle, x_j: node->cm */
388          r = sqrt((particle->pos_x - (*node)->cm_x)*(particle->pos_x - (*node)->cm_x) + ...
              (particle->pos_y - (*node)->cm_y)*(particle->pos_y - (*node)->cm_y));
389          F[0] += (*node)->mass_tot*(particle->pos_x - ...
              (*node)->cm_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
390          F[1] += (*node)->mass_tot*(particle->pos_y - ...
              (*node)->cm_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
391      }
392      /* Many body problem */
393      else
394      {
395          calculateForce(&(*node)->top_left, particle, F, theta_max);
396          calculateForce(&(*node)->top_right, particle, F, theta_max);
397          calculateForce(&(*node)->bottom_left, particle, F, theta_max);
398          calculateForce(&(*node)->bottom_right, particle, F, theta_max);
399      }
400  }
401  }

```

Appendix B

Listing 2: *Optimized galsim.c from Assignment 3*

```
1  #include <stdlib.h>
2  #include <math.h>
3  #include <stdio.h>
4  #include <sys/time.h>
5
6  typedef struct particle
7  {
8      double pos_x;
9      double pos_y;
10     double mass;
11     double vel_x;
12     double vel_y;
13     double brightness;
14 }particle_t;
15
16 int getInput(const char* __restrict filename, particle_t* __restrict data, const int N);
17 int setOutput(const char* __restrict filename, particle_t* __restrict data, const int N);
18 void calculateNewPositions(particle_t* data, const int N, const int nsteps, const double dt);
19
20 static double get_wall_seconds() {
21     struct timeval tv;
22     gettimeofday(&tv, NULL);
23     double seconds = tv.tv_sec + (double)tv.tv_usec / 1000000;
24     return seconds;
25 }
26
27 int main(int argc, char* argv[])
28 {
29     double time1 = get_wall_seconds();
30     if (argc != 6)
31     {
32         printf("Check your input..\n");
33         return -1;
34     }
35     const int N = atoi(argv[1]);
36     const char* filename = argv[2];
37     const int nsteps = atoi(argv[3]);
38     const double dt = atof(argv[4]);
39     //const int graphics = atoi(argv[5]);
40     const char* filename_out = "result.gal";
41     particle_t* data = (particle_t*)malloc(N*sizeof(particle_t));
42
43     int err = getInput(filename, data, N);
44     if(err == -1)
45     {
46         return -1;
47     }
48     calculateNewPositions(data, N, nsteps, dt);
49
50     err = setOutput(filename_out, data, N);
51     if(err == -1)
52     {
53         return -1;
54     }
55
56     free(data);
57     printf("galsim main took %7.3f wall seconds.\n", get_wall_seconds()-time1);
58     return 0;
59 }
```

```

60
61 int getInput(const char* __restrict filename, particle_t* __restrict data, const int N)
62 {
63     FILE *stream_in;
64     stream_in = fopen(filename, "rb");
65     double* arr = (double*)malloc(N*6*sizeof(double));
66
67     if(stream_in == NULL)
68     {
69         printf("Error: unable to open file: %s\n", filename);
70         return -1;
71     }
72     size_t input_size = N*6*sizeof(double);
73
74     /* Read input to data array */
75     size_t items_read = fread(arr, sizeof(char), input_size, stream_in);
76     int j = 0;
77     for (int i = 0; i < N*6; i += 6)
78     {
79         data[j].pos_x = arr[i];
80         data[j].pos_y = arr[i+1];
81         data[j].mass = arr[i+2];
82         data[j].vel_x = arr[i+3];
83         data[j].vel_y = arr[i+4];
84         data[j].brightness = arr[i+5];
85         j++;
86     }
87
88     if (items_read != input_size)
89     {
90         printf("Error reading the input file.\n");
91         return -1;
92     }
93
94     fclose(stream_in);
95     free(arr);
96     return 1;
97 }
98
99 int setOutput(const char* __restrict filename, particle_t* __restrict data, const int N)
100 {
101     FILE *stream_out;
102     stream_out = fopen(filename, "wb");
103     if(stream_out == NULL)
104     {
105         printf("Error: unable to open file: %s\n", filename);
106         return -1;
107     }
108     fwrite(data, sizeof(particle_t), N, stream_out);
109     fclose(stream_out);
110     return 0;
111 }
112
113 void calculateNewPositions(particle_t* particles, const int N, const int nsteps, const ...
    double dt)
114 {
115     double Fx, Fy, r, ax, ay;
116     const double G = 100/(double) (N);
117     const double epsilon = 0.001;
118     particle_t *particles_new = (particle_t*)malloc(N*sizeof(particle_t));
119     for (int n = 0; n < nsteps; n++)
120     {
121         for(int i = 0; i < N; i++)
122         {
123             Fx = 0;

```

```

124     Fy = 0;
125     for(int j = 0; j < i; j++)
126     {
127         r = sqrt((particles[i].pos_x - particles[j].pos_x)*(particles[i].pos_x - ...
            particles[j].pos_x) + (particles[i].pos_y - ...
            particles[j].pos_y)*(particles[i].pos_y - particles[j].pos_y));
128         Fx += particles[j].mass*(particles[i].pos_x - ...
            particles[j].pos_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
129         Fy += particles[j].mass*(particles[i].pos_y - ...
            particles[j].pos_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
130     }
131     for(int j = i+1; j < N; j++)
132     {
133         r = sqrt((particles[i].pos_x - particles[j].pos_x)*(particles[i].pos_x - ...
            particles[j].pos_x) + (particles[i].pos_y - ...
            particles[j].pos_y)*(particles[i].pos_y - particles[j].pos_y));
134         Fx += particles[j].mass*(particles[i].pos_x - ...
            particles[j].pos_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
135         Fy += particles[j].mass*(particles[i].pos_y - ...
            particles[j].pos_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
136     }
137     /* Forces */
138     Fx *= -G*particles[i].mass;
139     Fy *= -G*particles[i].mass;
140     /* Acceleration */
141     ax = Fx/(particles[i].mass);
142     ay = Fy/(particles[i].mass);
143     /* Updating velocities */
144     particles_new[i].vel_x = particles[i].vel_x + dt*ax;
145     particles_new[i].vel_y = particles[i].vel_y + dt*ay;
146     /* Updating positions */
147     particles_new[i].pos_x = particles[i].pos_x + dt* particles_new[i].vel_x;
148     particles_new[i].pos_y = particles[i].pos_y + dt* particles_new[i].vel_y;
149 }
150 for (int i = 0; i < N; i++)
151 {
152     particles[i].pos_x = particles_new[i].pos_x;
153     particles[i].pos_y = particles_new[i].pos_y;
154     particles[i].vel_x = particles_new[i].vel_x;
155     particles[i].vel_y = particles_new[i].vel_y;
156 }
157 }
158 free(particles_new);
159 }

```

Appendix C

Listing 3: *Parallelized galsim.c from Assignment 3*

```
1  #include <stdlib.h>
2  #include <math.h>
3  #include <stdio.h>
4  #include <sys/time.h>
5  #include <pthread.h>
6
7  typedef struct particle
8  {
9      double pos_x;
10     double pos_y;
11     double mass;
12     double vel_x;
13     double vel_y;
14     double brightness;
15 }particle_t;
16
17 typedef struct dataForThread
18 {
19     int start;
20     int stop;
21     int N;
22     double dt;
23     pthread_t thread;
24 }dataForThread_t;
25
26 /* Global arrays */
27 particle_t* particles;
28 particle_t* particles_new;
29 int rest;
30
31 int getInput(const char* __restrict filename, const int N);
32 int setOutput(const char* __restrict filename, const int N);
33 void calculateNewPositions(const int N, const int nsteps, const double dt, int ...
    num_of_threads);
34 void* thread_func(void* arg);
35
36 static double get_wall_seconds() {
37     struct timeval tv;
38     gettimeofday(&tv, NULL);
39     double seconds = tv.tv_sec + (double)tv.tv_usec / 1000000;
40     return seconds;
41 }
42
43 int main(int argc, char* argv[])
44 {
45     double time1 = get_wall_seconds();
46     if (argc != 8)
47     {
48         printf("Expected input: ./galsim_A3 N filename nsteps Δ_t theta_max graphics ...
            n_threads\n");
49         return -1;
50     }
51     const int N = atoi(argv[1]);
52     const char* filename = argv[2];
53     const int nsteps = atoi(argv[3]);
54     const double dt = atof(argv[4]);
55     //const double theta_max = atof(argv[5]);
56     //const int graphics = atoi(argv[6]);
57     const int num_of_threads = atoi(argv[7]);
```

```

58
59     rest = N%num_of_threads;
60
61     const char* filename_out = "result.gal";
62     particles = (particle_t*)malloc(N*sizeof(particle_t));
63
64     int err = getInput(filename, N);
65     if(err == -1)
66     {
67         return -1;
68     }
69
70     calculateNewPositions(N, nsteps, dt, num_of_threads);
71
72     err = setOutput(filename_out, N);
73     if(err == -1)
74     {
75         return -1;
76     }
77
78     free(particles);
79     printf("galsim main took %7.3f wall seconds.\n", get_wall_seconds()-time1);
80     return 0;
81 }
82
83 int getInput(const char* __restrict filename, const int N)
84 {
85     FILE *stream_in;
86     stream_in = fopen(filename, "rb");
87     double* arr = (double*)malloc(N*6*sizeof(double));
88
89     if(stream_in == NULL)
90     {
91         printf("Error: unable to open file: %s\n", filename);
92         return -1;
93     }
94     size_t input_size = N*6*sizeof(double);
95
96     /* Read input to data array */
97     size_t items_read = fread(arr, sizeof(char), input_size, stream_in);
98     int j = 0;
99     for (int i = 0; i < N*6; i += 6)
100     {
101         particles[j].pos_x = arr[i];
102         particles[j].pos_y = arr[i+1];
103         particles[j].mass = arr[i+2];
104         particles[j].vel_x = arr[i+3];
105         particles[j].vel_y = arr[i+4];
106         particles[j].brightness = arr[i+5];
107         j++;
108     }
109
110     if (items_read != input_size)
111     {
112         printf("Error reading the input file.\n");
113         return -1;
114     }
115
116     fclose(stream_in);
117     free(arr);
118     return 1;
119 }
120
121 int setOutput(const char* __restrict filename, const int N)
122 {

```



```

123 FILE *stream_out;
124 stream_out = fopen(filename, "wb");
125 if (stream_out == NULL)
126 {
127     printf("Error: unable to open file: %s\n", filename);
128     return -1;
129 }
130 fwrite(particles, sizeof(particle_t), N, stream_out);
131 fclose(stream_out);
132 return 0;
133 }
134
135 void calculateNewPositions(const int N, const int nsteps, const double dt, int num_of_threads)
136 {
137     dataForThread_t threads[num_of_threads];
138     particles_new = (particle_t*)malloc(N*sizeof(particle_t));
139     for (int n = 0; n < nsteps; n++)
140     {
141         int work_size = N/num_of_threads; // number of particles for every thread
142         /* Create threads */
143         for(int i = 0; i < num_of_threads-1; i++){
144             threads[i].start = (i+1) * work_size - work_size;
145             threads[i].stop = threads[i].start + work_size;
146             threads[i].N = N;
147             threads[i].dt = dt;
148
149             pthread_create(&(threads[i].thread), NULL, thread_func, &threads[i]);
150         }
151         /* Create last thread (with possibly different work_size) */
152         threads[num_of_threads-1].start = num_of_threads * work_size - work_size;
153         threads[num_of_threads-1].stop = threads[num_of_threads-1].start + work_size + rest;
154         threads[num_of_threads-1].N = N;
155         threads[num_of_threads-1].dt = dt;
156         pthread_create(&(threads[num_of_threads-1].thread), NULL, thread_func, ...
157             &threads[num_of_threads-1]);
158
159         /* Join threads */
160         for(int i = 0; i < num_of_threads; i++){
161             pthread_join(threads[i].thread, NULL);
162         }
163
164         for (int i = 0; i < N; i++)
165         {
166             particles[i].pos_x = particles_new[i].pos_x;
167             particles[i].pos_y = particles_new[i].pos_y;
168             particles[i].vel_x = particles_new[i].vel_x;
169             particles[i].vel_y = particles_new[i].vel_y;
170         }
171         free(particles_new);
172     }
173 }
174
175 void* thread_func(void* arg){
176     dataForThread_t* info = (dataForThread_t *) arg;
177
178     double Fx, Fy, r, ax, ay;
179     const double G = 100/(double) (info->N);
180     const double epsilon = 0.001;
181     /* Compute new position of particle */
182     for(int i = info->start; i<info->stop; i++){ // for specific particles
183         Fx = 0;
184         Fy = 0;
185         for(int j = 0; j<i; j++){ // compare with every other particle
186             /* Distance between particles */

```

```

186     r = sqrt((particles[i].pos_x - particles[j].pos_x)*(particles[i].pos_x - ...
        particles[j].pos_x) + (particles[i].pos_y - ...
        particles[j].pos_y)*(particles[i].pos_y - particles[j].pos_y));
187
188     /* Forces*/
189     Fx += particles[j].mass*(particles[i].pos_x - ...
        particles[j].pos_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
190     Fy += particles[j].mass*(particles[i].pos_y - ...
        particles[j].pos_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
191 }
192 for(int j = i+1; j < info->N; j++){
193     /* Distance between particles */
194     r = sqrt((particles[i].pos_x - particles[j].pos_x)*(particles[i].pos_x - ...
        particles[j].pos_x) + (particles[i].pos_y - ...
        particles[j].pos_y)*(particles[i].pos_y - particles[j].pos_y));
195
196     /* Forces*/
197     Fx += particles[j].mass*(particles[i].pos_x - ...
        particles[j].pos_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
198     Fy += particles[j].mass*(particles[i].pos_y - ...
        particles[j].pos_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
199 }
200 /* Forces */
201 Fx -= -G * particles[i].mass;
202 Fy -= -G * particles[i].mass;
203 /* Acceleration */
204 ax = Fx/(particles[i].mass);
205 ay = Fy/(particles[i].mass);
206 /* Updating velocities */
207 particles_new[i].vel_x = particles[i].vel_x + info->dt * ax;
208 particles_new[i].vel_y = particles[i].vel_y + info->dt * ay;
209 /* Updating positions */
210 particles_new[i].pos_x = particles[i].pos_x + (double)info->dt * particles_new[i].vel_x;
211 particles_new[i].pos_y = particles[i].pos_y + (double)info->dt * particles_new[i].vel_y;
212 }
213 return NULL;
214 }

```

Listing 4: *Parallelized galsim.c from Assignment 4*

```

1  #include <stdlib.h>
2  #include <math.h>
3  #include <stdio.h>
4  #include <sys/time.h>
5  #include <pthread.h>
6
7  typedef struct particle
8  {
9      double pos_x;
10     double pos_y;
11     double mass;
12     double vel_x;
13     double vel_y;
14     double brightness;
15 }__attribute__((packed__)) particle_t;
16
17 typedef struct quadrant
18 {
19     /* Position of specific quadrant in the whole space */
20     double origin_x;
21     double origin_y;
22
23     double width;

```

```

24     double height;
25 } __attribute__((__packed__)) quadrant_t;
26
27 typedef struct quadtree
28 {
29     /* center of mass */
30     double mass_tot;
31     double cm_x;
32     double cm_y;
33     particle_t *particle;
34     quadrant_t *quadrant;
35     /* tree nodes */
36     struct quadtree *top_left;
37     struct quadtree *top_right;
38     struct quadtree *bottom_left;
39     struct quadtree *bottom_right;
40 } quadtree_t;
41
42 typedef struct dataForThread
43 {
44     int start;
45     int stop;
46     int N;
47     double dt;
48     double theta_max;
49     pthread_t thread;
50     quadtree_t** root;
51 }dataForThread_t;
52
53 /* Global arrays */
54 particle_t* particles;
55 particle_t* particles_new;
56 int rest;
57
58 int getInput(const char* __restrict filename, quadtree_t** __restrict root, const int N);
59 int insert(quadtree_t** __restrict node, particle_t* __restrict particle, quadrant_t* ...
    __restrict quadrant);
60 void insertQuadNode(quadtree_t** node, particle_t* __restrict particle);
61 void deleteQuadtree(quadtree_t **node);
62 int setOutput(const char* __restrict filename, const int N);
63 int calculateNewPositions(quadtree_t** __restrict root, const int N, const int nsteps, ...
    const double dt, const double theta_max, int num_of_threads);
64 void calculateForce(quadtree_t** __restrict node, particle_t* __restrict particle, double* ...
    __restrict F, const double theta_max);
65 void* thread_func(void* arg);
66
67 static double get_wall_seconds() {
68     struct timeval tv;
69     gettimeofday(&tv, NULL);
70     double seconds = tv.tv_sec + (double)tv.tv_usec / 1000000;
71     return seconds;
72 }
73
74 int main(int argc, char* argv[])
75 {
76     double time1 = get_wall_seconds();
77     if (argc != 8)
78     {
79         printf("Expected input: ./galsim_A4 N filename nsteps Δ_t theta_max graphics ...
            n_threads\n");
80         return -1;
81     }
82     const int N = atoi(argv[1]);
83     const char* filename = argv[2];
84     const int nsteps = atoi(argv[3]);

```

```

85     const double dt = atof(argv[4]);
86     const double theta_max = atof(argv[5]); // theta_max = 0.25
87     //const int graphics = atoi(argv[6]);
88     const int num_of_threads = atoi(argv[7]);
89
90     rest = N%num_of_threads;
91
92     const char* filename_out = "result.gal";
93
94     particles = (particle_t*)malloc(N*sizeof(particle_t));
95     quadtree_t* root = NULL;
96
97     int err = getInput(filename, &root, N);
98     if(err == -1)
99     {
100         deleteQuadtree(&root);
101         free(particles);
102         return -1;
103     }
104
105     err = calculateNewPositions(&root, N, nsteps, dt, theta_max, num_of_threads);
106     if(err == -1)
107     {
108         deleteQuadtree(&root);
109         free(particles);
110         return -1;
111     }
112
113     err = setOutput(filename_out, N);
114     if(err == -1)
115     {
116         deleteQuadtree(&root);
117         free(particles);
118         return -1;
119     }
120
121     deleteQuadtree(&root);
122     free(particles);
123     printf("galsim main took %7.3f wall seconds.\n", get_wall_seconds()-timel);
124     return 0;
125 }
126
127 int getInput(const char* __restrict filename, quadtree_t** __restrict root, const int N)
128 {
129     FILE *stream_in;
130     stream_in = fopen(filename, "rb");
131     double* arr = (double*)malloc(N*6*sizeof(double));
132     quadrant_t *quadrant = (quadrant_t*)malloc(sizeof(particle_t));
133
134     if(stream_in == NULL)
135     {
136         printf("Error: unable to open file: %s\n", filename);
137         fclose(stream_in);
138         free(arr);
139         free(quadrant);
140         return -1;
141     }
142     size_t input_size = N*6*sizeof(double);
143
144     /* Read input to data array */
145     size_t items_read = fread(arr, sizeof(char), input_size, stream_in);
146     int j = 0;
147     for (int i = 0; i < N*6; i += 6)
148     {
149         particles[j].pos_x = arr[i];

```

```

150     particles[j].pos_x = arr[i];
151     particles[j].pos_y = arr[i+1];
152     particles[j].mass = arr[i+2];
153     particles[j].vel_x = arr[i+3];
154     particles[j].vel_y = arr[i+4];
155     particles[j].brightness = arr[i+5];
156
157     quadrant->width = 1;
158     quadrant->height = 1;
159     quadrant->origin_x = 0;
160     quadrant->origin_y = 0;
161
162     insert(root, &particles[j], quadrant);
163     j++;
164 }
165
166 if (items_read != input_size)
167 {
168     printf("Error reading the input file.\n");
169     fclose(stream_in);
170     free(arr);
171     free(quadrant);
172     return -1;
173 }
174
175 free(quadrant);
176 fclose(stream_in);
177 free(arr);
178 return 1;
179 }
180
181 int insert(quadtrees_t** node, particle_t* __restrict particle, quadrant_t* __restrict ...
182     quadrant)
183 {
184     if (*node == NULL)
185     {
186         quadtrees_t * new_node = (quadtrees_t*)malloc(sizeof(quadtrees_t));
187         quadrant_t * new_quadrant = (quadrant_t*)malloc(sizeof(quadrant_t));
188         new_node->particle = particle;
189         new_node->mass_tot = particle->mass;
190         new_node->cm_x = particle->pos_x;
191         new_node->cm_y = particle->pos_y;
192
193         new_quadrant->width = quadrant->width;
194         new_quadrant->height = quadrant->height;
195         new_quadrant->origin_x = quadrant->origin_x;
196         new_quadrant->origin_y = quadrant->origin_y;
197         new_node->quadrant = new_quadrant;
198
199         new_node->top_left = NULL;
200         new_node->top_right = NULL;
201         new_node->bottom_left = NULL;
202         new_node->bottom_right = NULL;
203         *node = new_node;
204     }
205     else if ((*node)->particle == NULL) // internal node ("stores" several particles)
206     {
207         /* updating center of mass */
208         (*node)->cm_x = ((*node)->cm_x*(*node)->mass_tot + ...
209             particle->pos_x*particle->mass)/((*node)->mass_tot + particle->mass);
210         (*node)->cm_y = ((*node)->cm_y*(*node)->mass_tot + ...
211             particle->pos_y*particle->mass)/((*node)->mass_tot + particle->mass);
212         (*node)->mass_tot += particle->mass;
213         insertQuadNode(&(*node), particle);
214     }
215 }

```

```

212 else // external node (node with only one particle)
213 {
214     if ((*node)->particle->pos_x == particle->pos_x && (*node)->particle->pos_y == ...
215         particle->pos_y) {
216         printf("Error: Particles at the same position \n");
217         return -1;
218     }
219     /* Making our node internal by removing the particle */
220     particle_t *particle_old = (*node)->particle;
221     (*node)->particle = NULL;
222     insertQuadNode(&(*node), particle);
223     insertQuadNode(&(*node), particle_old);
224
225     /* updating center of mass */
226     (*node)->cm_x = ((*node)->cm_x*(*node)->mass_tot + ...
227         particle->pos_x*particle->mass)/((*node)->mass_tot + particle->mass);
228     (*node)->cm_y = ((*node)->cm_y*(*node)->mass_tot + ...
229         particle->pos_y*particle->mass)/((*node)->mass_tot + particle->mass);
230     (*node)->mass_tot += particle->mass;
231 }
232 return 0;
233 }
234
235 void insertQuadNode(quadtrees_t** __restrict node, particle_t* __restrict particle)
236 {
237     quadrant_t *quadrant_new = (quadrant_t*)malloc(sizeof(quadrant_t));
238     quadrant_new->width = (*node)->quadrant->width/2.0;
239     quadrant_new->height = (*node)->quadrant->height/2.0;
240
241     if (particle->pos_x < ((*node)->quadrant->origin_x + quadrant_new->width)) // left quadrant
242     {
243         if (particle->pos_y < ((*node)->quadrant->origin_y + quadrant_new->height)) // ...
244             bottom left quadrant
245         {
246             quadrant_new->origin_y = (*node)->quadrant->origin_y;
247             quadrant_new->origin_x = (*node)->quadrant->origin_x;
248             insert(&(*node)->bottom_left, particle, quadrant_new);
249         }
250         else // top left quadrant
251         {
252             quadrant_new->origin_y = (*node)->quadrant->origin_y + quadrant_new->height;
253             quadrant_new->origin_x = (*node)->quadrant->origin_x;
254             insert(&(*node)->top_left, particle, quadrant_new);
255         }
256     }
257     else { // right quadrant
258         if (particle->pos_y < ((*node)->quadrant->origin_y + quadrant_new->height)) // ...
259             bottom right quadrant
260         {
261             quadrant_new->origin_x = (*node)->quadrant->origin_x + quadrant_new->width;
262             quadrant_new->origin_y = (*node)->quadrant->origin_y;
263             insert(&(*node)->bottom_right, particle, quadrant_new);
264         }
265         else // top right quadrant
266         {
267             quadrant_new->origin_x = (*node)->quadrant->origin_x + quadrant_new->width;
268             quadrant_new->origin_y = (*node)->quadrant->origin_y + quadrant_new->height;
269             insert(&(*node)->top_right, particle, quadrant_new);
270         }
271     }
272     free(quadrant_new);
273 }
274
275 void deleteQuadtree(quadtrees_t **node)
276 {

```

```

272     if(*node == NULL){ // if empty node
273         return;
274     }
275
276     deleteQuadtree(&(*node)->top_left);
277     deleteQuadtree(&(*node)->top_right);
278     deleteQuadtree(&(*node)->bottom_left);
279     deleteQuadtree(&(*node)->bottom_right);
280
281     free((*node)->quadrant);
282     free(*node);
283     *node = NULL;
284 }
285
286 int setOutput(const char* __restrict filename, const int N)
287 {
288     FILE *stream_out;
289     stream_out = fopen(filename, "wb");
290     if(stream_out == NULL)
291     {
292         printf("Error: unable to open file: %s\n", filename);
293         fclose(stream_out);
294         return -1;
295     }
296     fwrite(particles, sizeof(particle_t), N, stream_out);
297     fclose(stream_out);
298     return 0;
299 }
300
301 int calculateNewPositions(quadtree_t** root, const int N, const int nsteps, const double ...
    dt, const double theta_max, int num_of_threads)
302 {
303     int err = 0;
304     particles_new = (particle_t*)malloc(N*sizeof(particle_t));
305     quadrant_t* quadrant = (quadrant_t*)malloc(sizeof(quadrant_t));
306
307     quadtree_t **root_new = (quadtree_t**)malloc(sizeof(quadtree_t));
308     *root_new = *root;
309
310     dataForThread_t threads[num_of_threads];
311     for (int n = 0; n < nsteps; n++)
312     {
313         int work_size = N/num_of_threads; // number of particles for every thread
314         /* Create threads */
315         for(int i = 0; i < num_of_threads-1; i++){
316             threads[i].start = (i+1) * work_size - work_size;
317             threads[i].stop = threads[i].start + work_size;
318             threads[i].N = N;
319             threads[i].dt = dt;
320             threads[i].theta_max = theta_max;
321             threads[i].root = root_new;
322
323             pthread_create(&(threads[i].thread), NULL, thread_func, &threads[i]);
324         }
325         /* Create last thread (with possibly different work_size) */
326         threads[num_of_threads-1].start = num_of_threads * work_size - work_size;
327         threads[num_of_threads-1].stop = threads[num_of_threads-1].start + work_size + rest;
328         threads[num_of_threads-1].N = N;
329         threads[num_of_threads-1].dt = dt;
330         threads[num_of_threads-1].theta_max = theta_max;
331         threads[num_of_threads-1].root = root_new;
332         pthread_create(&(threads[num_of_threads-1].thread), NULL, thread_func, ...
            &threads[num_of_threads-1]);
333
334         /* Join threads */

```

```

335     for(int i = 0; i < num_of_threads; i++){
336         pthread_join(threads[i].thread, NULL);
337     }
338
339     /* Check that particles are within interval */
340     for(int i = 0; i < N; i++){
341         if(particles_new[i].pos_x > 1 || particles_new[i].pos_x < 0 || particles_new[i].pos_y ...
           > 1 || particles_new[i].pos_y < 0){
342             printf("Error: Particle outside interval\n");
343             free(particles_new);
344             free(root_new);
345             free(quadrant);
346             return -1;
347         }
348     }
349
350     deleteQuadtrees(root_new); // reset tree
351
352     /* Update particle positions and refill tree */
353     for (int i = 0; i < N; i++)
354     {
355         particles[i].pos_x = particles_new[i].pos_x;
356         particles[i].pos_y = particles_new[i].pos_y;
357         particles[i].vel_x = particles_new[i].vel_x;
358         particles[i].vel_y = particles_new[i].vel_y;
359
360         quadrant->width = 1;
361         quadrant->height = 1;
362         quadrant->origin_x = 0;
363         quadrant->origin_y = 0;
364         err += insert(root_new, &particles[i], quadrant);
365     }
366     if(err!=0)
367     {
368         free(particles_new);
369         free(root_new);
370         free(quadrant);
371         return -1;
372     }
373 }
374 *root = *root_new;
375 free(particles_new);
376 free(root_new);
377 free(quadrant);
378 return 0;
379 }
380
381 void calculateForce(quadtrees_t* __restrict node, particle_t* __restrict particle, double* ...
           __restrict F, const double theta_max)
382 {
383     double r, theta;
384     const double epsilon = 0.001;
385     if ((*node) == NULL)
386     {
387         return;
388     }
389     /* External node */
390     else if ((*node)->particle != NULL)
391     {
392         /* Same particle */
393         if ((fabs((*node)->particle->pos_x - particle->pos_x) < 0.0000001) && ...
           (fabs((*node)->particle->pos_y - particle->pos_y) < 0.0000001))
394         {
395             return;
396         }

```



```

397     r = sqrt((particle->pos_x - (*node)->particle->pos_x)*(particle->pos_x - ...
        (*node)->particle->pos_x) + (particle->pos_y - ...
        (*node)->particle->pos_y)*(particle->pos_y - (*node)->particle->pos_y));
398     F[0] += (*node)->particle->mass*(particle->pos_x - ...
        (*node)->particle->pos_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
399     F[1] += (*node)->particle->mass*(particle->pos_y - ...
        (*node)->particle->pos_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
400 }
401 /* Internal node */
402 else
403 {
404     /* Calculate ratio */
405     theta = (*node)->quadrant->width/sqrt((particle->pos_x - ((*node)->quadrant->origin_x ...
        + (*node)->quadrant->width/2))* (particle->pos_x - ((*node)->quadrant->origin_x + ...
        (*node)->quadrant->width/2)) + (particle->pos_y - ((*node)->quadrant->origin_y + ...
        (*node)->quadrant->height/2))* (particle->pos_y - ((*node)->quadrant->origin_y + ...
        (*node)->quadrant->height/2)));
406
407     /* Single body problem */
408     if (theta ≤ theta_max)
409     {
410         /* x_i: particle, x_j: node->cm */
411         r = sqrt((particle->pos_x - (*node)->cm_x)*(particle->pos_x - (*node)->cm_x) + ...
            (particle->pos_y - (*node)->cm_y)*(particle->pos_y - (*node)->cm_y));
412         F[0] += (*node)->mass_tot*(particle->pos_x - ...
            (*node)->cm_x)/((r+epsilon)*(r+epsilon)*(r+epsilon));
413         F[1] += (*node)->mass_tot*(particle->pos_y - ...
            (*node)->cm_y)/((r+epsilon)*(r+epsilon)*(r+epsilon));
414     }
415     /* Many body problem */
416     else
417     {
418         calculateForce(&(*node)->top_left, particle, F, theta_max);
419         calculateForce(&(*node)->top_right, particle, F, theta_max);
420         calculateForce(&(*node)->bottom_left, particle, F, theta_max);
421         calculateForce(&(*node)->bottom_right, particle, F, theta_max);
422     }
423 }
424 }
425
426 void* thread_func(void* arg){
427     dataForThread_t* info = (dataForThread_t *) arg;
428
429     double* F = (double*)malloc(2*sizeof(double));
430     double ax, ay;
431     const double G = 100/(double)(info->N);
432
433     /* Compute new position of particle */
434     for(int i = info->start; i<info->stop; i++){ // for specific particles
435         F[0] = 0;
436         F[1] = 0;
437
438         calculateForce(info->root, &particles[i], F, info->theta_max);
439
440         F[0] = -G*particles[i].mass*F[0];
441         F[1] = -G*particles[i].mass*F[1];
442
443         /* Acceleration */
444         ax = F[0]/(particles[i].mass);
445         ay = F[1]/(particles[i].mass);
446         /* Updating velocities */
447         particles_new[i].vel_x = particles[i].vel_x + info->dt * ax;
448         particles_new[i].vel_y = particles[i].vel_y + info->dt * ay;
449         /* Updating positions */
450         particles_new[i].pos_x = particles[i].pos_x + info->dt * particles_new[i].vel_x;

```

```
451     particles_new[i].pos_y = particles[i].pos_y + info->dt * particles_new[i].vel_y;
452 }
453 free(F);
454 return NULL;
455 }
```